Question #1 *Topic 2*

HOTSPOT -

You are developing an application that will use the Computer Vision client library. The application has the following code.

```
public async TaskAnalyzeImage(ComputerVisionClient client, string localImage)
{
        List<VisualFeatureTypes> features = new List<VisualFeatureTypes>()
                {
                        VisualFeatureTypes.Description,
                        VisualFeatureTypes.Tags,
                };
        using (Stream imageStream = File.OpenRead(localImage))
        {
                try
                {
                  ImageAnalysis results = await client.AnalyzeImageInStreamAsync(imageStream, features);

                  foreach (var caption in results.Description.Captions)
                  {
                       Console.WriteLine($"{caption.Text} with confidence {caption.Confidence}");
                  }

                  foreach (var tag in results.Tags)
                  {
                       Console.WriteLine($"{tag.Name} {tag.Confidence}");
                  }
                }
                catch (Exception ex)
                {
                        Console.WriteLine(ex.Message);
                }
        }
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
The code will perform face recognition.	○	○
The code will list tags and their associated confidence.	○	○
The code will read a file from the local file system.	○	○

Correct Answer:

Answer Area

Statements	Yes	No
The code will perform face recognition.	○	◉
The code will list tags and their associated confidence.	◉	○
The code will read a file from the local file system.	○	◉

Box 1: No -

Box 2: Yes -

The ComputerVision.analyzeImageInStreamAsync operation extracts a rich set of visual features based on the image content.

Box 3: No -

Images will be read from a stream.

Reference:

https://docs.microsoft.com/en-us/java/api/com.microsoft.azure.cognitiveservices.vision.computervision.computervision.analyzeimageinstreamasync

Question #2 *Topic 2*

You are developing a method that uses the Computer Vision client library. The method will perform optical character recognition (OCR) in images. The method has the following code.

```
public static async Task ReadFileUrl(ComputerVisionClient client, string urlFile)
{
        const int numberOfCharsInOperationId = 36;

        var txtHeaders = await client.ReadAsync(urlFile, language: "en");

        string opLocation = textHeaders.OperationLocation;
        string operationId = opLocation.Substring(opLocation.Length -
numberOfCharsInOperationId);

        ReadOperationResult results;

        results = await client.GetReadResultAsync(Guid.Parse(operationId));

        var textUrlFileResults = results.AnalyzeResult.ReadResults;
        foreach (ReadResult page in textUrlFileResults)
        {
                foreach (Line line in page.Lines)
                {
                        Console.WriteLine(line.Text);
                }
        }
}
```

During testing, you discover that the call to the GetReadResultAsync method occurs before the read operation is complete.

You need to prevent the GetReadResultAsync method from proceeding until the read operation is complete.

Which two actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

A. Remove the Guid.Parse(operationId) parameter.

B. Add code to verify the results.Status value. **Most Voted**

C. Add code to verify the status of the txtHeaders.Status value.

D. Wrap the call to GetReadResultAsync within a loop that contains a delay. **Most Voted**

Correct Answer: *BD*

Community vote distribution

BD (100%)

Question #3 *Topic 2*

HOTSPOT -

You have a Computer Vision resource named contoso1 that is hosted in the West US Azure region.

You need to use contoso1 to make a different size of a product photo by using the smart cropping feature.

How should you complete the API URL? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
curl -H "Ocp-Apim-Subscription-Key: xxx" /
-o "sample.png" -H "Content-Type: application/json" /
[ ▼ ] /vision/v3.1/ [ ▼ ] ?width=100&height=100&smartCropping=true" /
-d "{\"url\":\"https://upload.litwareinc.org/litware/bicycle.jpg\"}"
```

- [▼]: "https://api.projectoxford.ai" / "https://contoso1.cognitiveservices.azure.com" / "https://westus.api.cognitive.microsoft.com"
- [▼]: areaOfInterest / detect / generateThumbnail

Correct Answer:

Answer Area

```
curl -H "Ocp-Apim-Subscription-Key: xxx" /
-o "sample.png" -H "Content-Type: application/json" /
[ ▼ ] /vision/v3.1/ [ areaOfInterest ▼ ] ?width=100&height=100&smartCropping=true" /
-d "{\"url\":\"https://upload.litwareinc.org/litware/bicycle.jpg\"}"
```

- [▼]: "https://westus.api.cognitive.microsoft.com"
- [areaOfInterest ▼]: areaOfInterest

Reference:

https://westus.dev.cognitive.microsoft.com/docs/services/computer-vision-v3-2/operations/56f91f2e778daf14a499f21b

https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-generating-thumbnails#examples

Question #4 *Topic 2*

DRAG DROP -

You are developing a webpage that will use the Azure Video Analyzer for Media (previously Video Indexer) service to display videos of internal company meetings.

You embed the Player widget and the Cognitive Insights widget into the page.

You need to configure the widgets to meet the following requirements:

- Ensure that users can search for keywords.
- Display the names and faces of people in the video.
- Show captions in the video in English (United States).

How should you complete the URL for each widget? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Correct Answer:

Reference:

https://docs.microsoft.com/en-us/azure/azure-video-analyzer/video-analyzer-for-media-docs/video-indexer-embed-widgets

Question #5 *Topic 2*

DRAG DROP -

You train a Custom Vision model to identify a company's products by using the Retail domain.

You plan to deploy the model as part of an app for Android phones.

You need to prepare the model for deployment.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

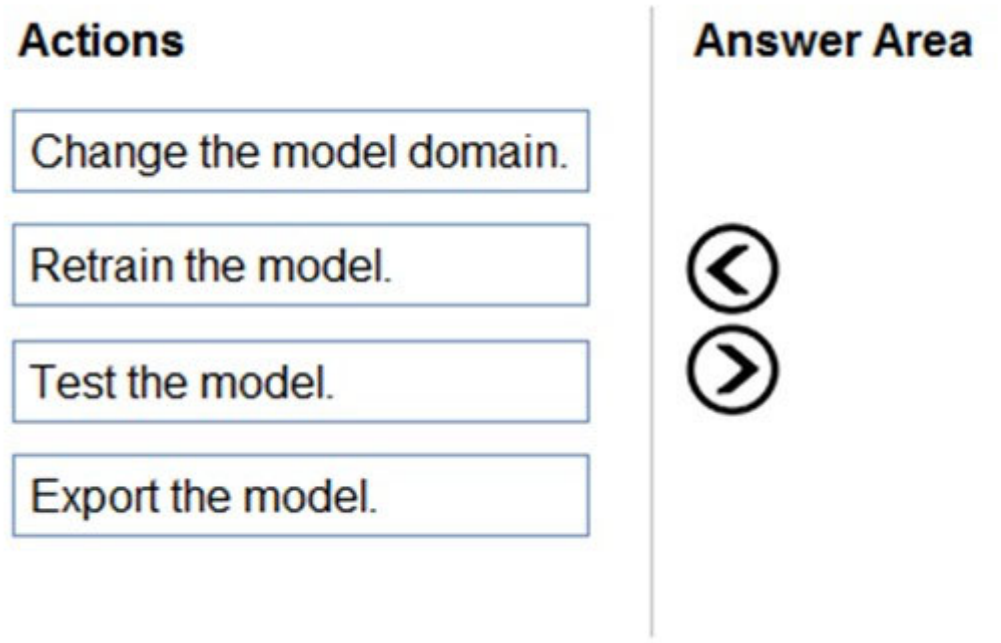

Correct Answer:

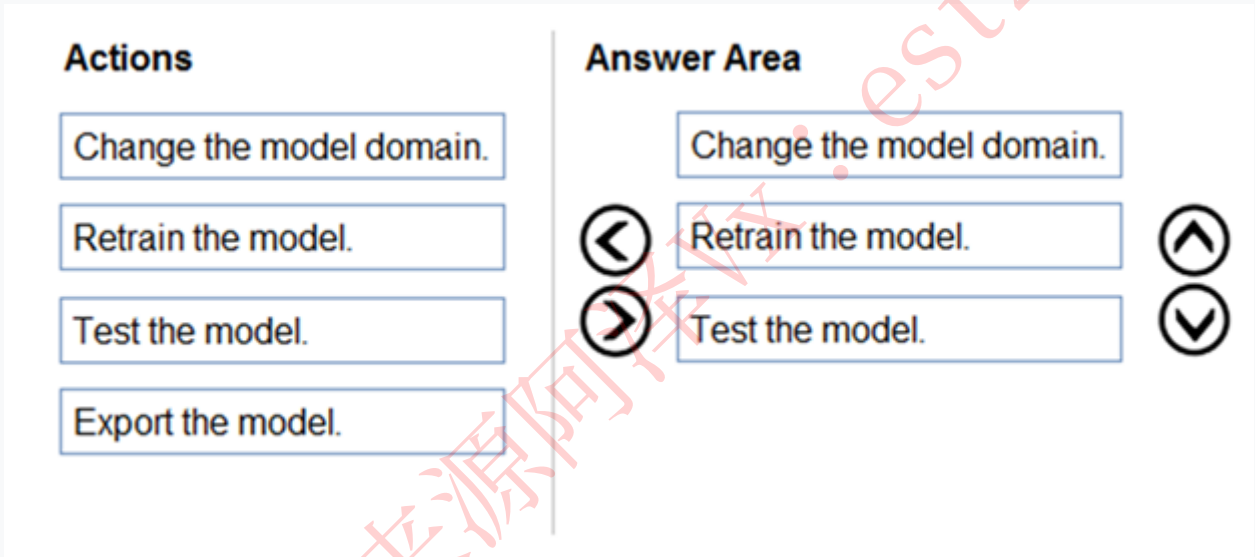

Reference:

https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/export-your-model

Question #6 *Topic 2*

HOTSPOT -

You are developing an application to recognize employees' faces by using the Face Recognition API. Images of the faces will be accessible from a URI endpoint.

The application has the following code.

```
def add_face(subscription_key, person_group_id, person_id, image_uri):
    headers = {
        'Content-Type': 'application/json',
        'Ocp-Apim-Subscription-Key': subscription_key
        }
    body = {
        'url': image_uri }
    conn = httplib.HTTPSConnection('westus.api.cognitive.microsoft.com')
    conn.request('POST',
f'/face/v1.0/persongroups/{person_group_id}/persons/{person_id}/persistedFaces', f'{body}', headers)
    response = conn.getresponse()
        response_data = response.read()
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
The code will add a face image to a person object in a person group.	○	○
The code will work for up to 10,000 people.	○	○
`add_face` can be called multiple times to add multiple face images to a person object.	○	○

Correct Answer:

Answer Area

Statements	Yes	No
The code will add a face image to a person object in a person group.	○	○
The code will work for up to 10,000 people.	○	○
`add_face` can be called multiple times to add multiple face images to a person object.	○	○

Reference:

https://docs.microsoft.com/en-us/azure/cognitive-services/face/face-api-how-to-topics/use-persondirectory

Question #7 *Topic 2*

DRAG DROP -

You have a Custom Vision resource named acvdev in a development environment.

You have a Custom Vision resource named acvprod in a production environment.

In acvdev, you build an object detection model named obj1 in a project named proj1.

You need to move obj1 to acvprod.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

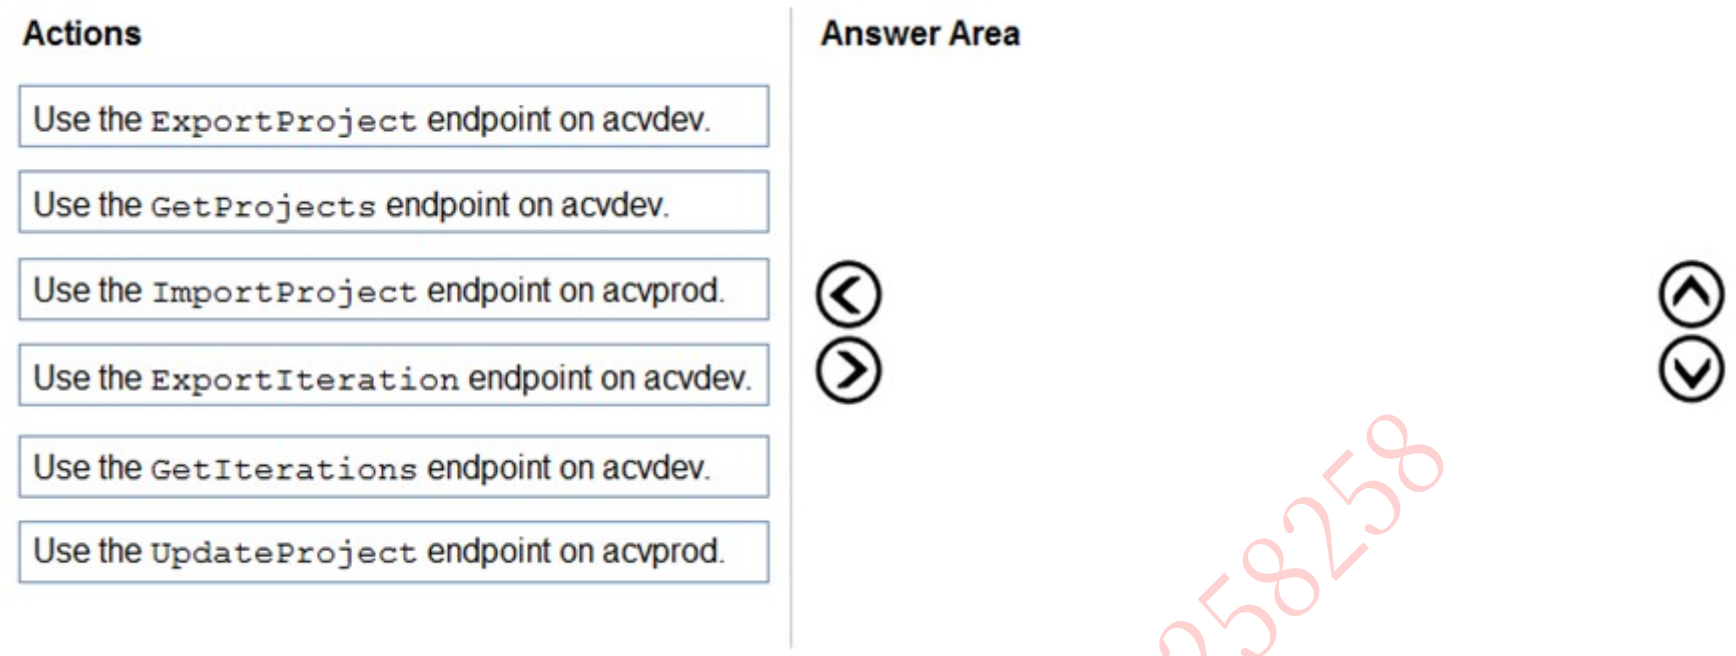

Correct Answer:

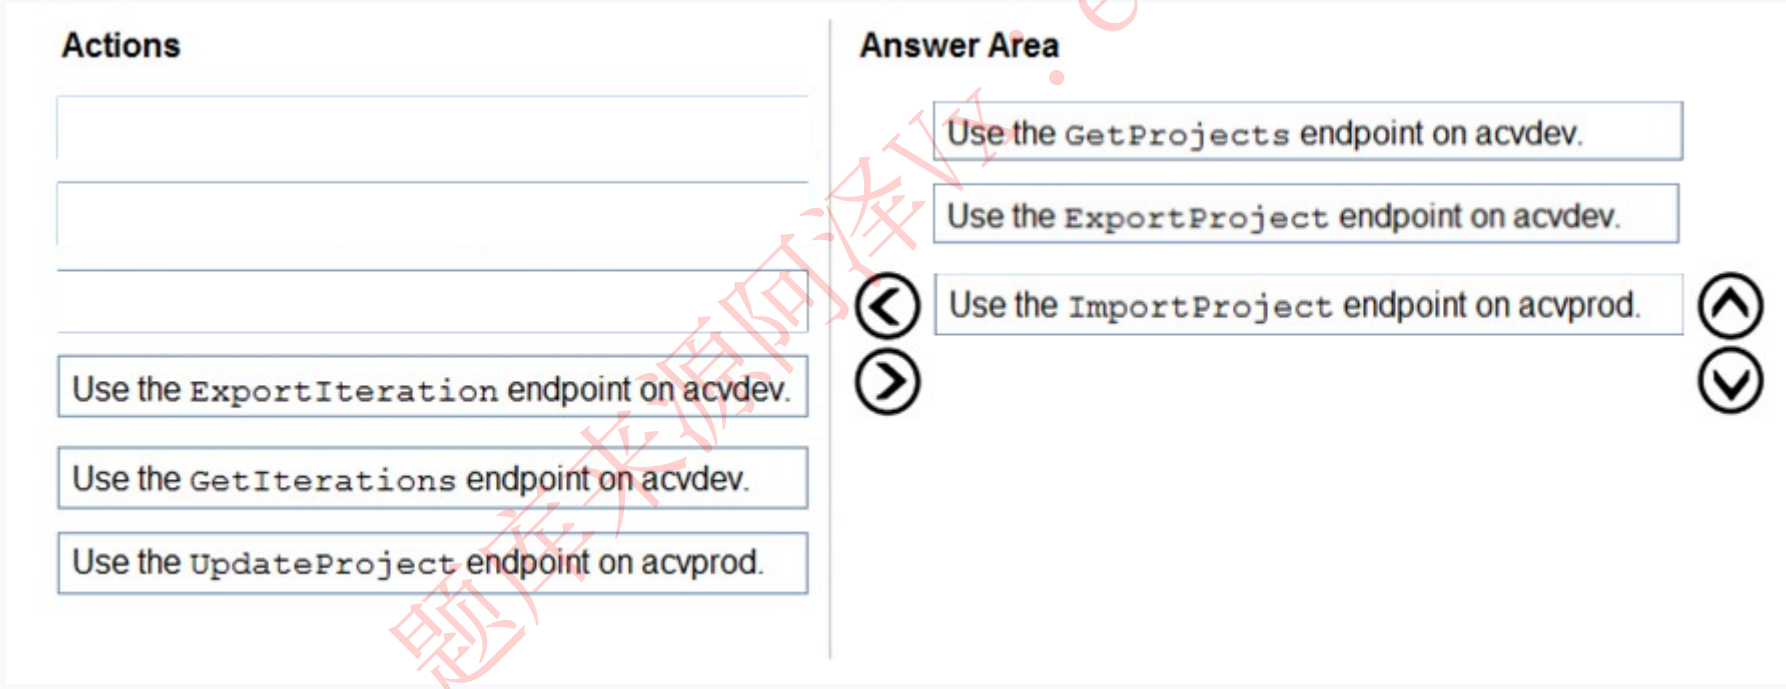

Reference:

https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/copy-move-projects

Question #8 *Topic 2*

DRAG DROP -

You are developing an application that will recognize faults in components produced on a factory production line. The components are specific to your business.

You need to use the Custom Vision API to help detect common faults.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

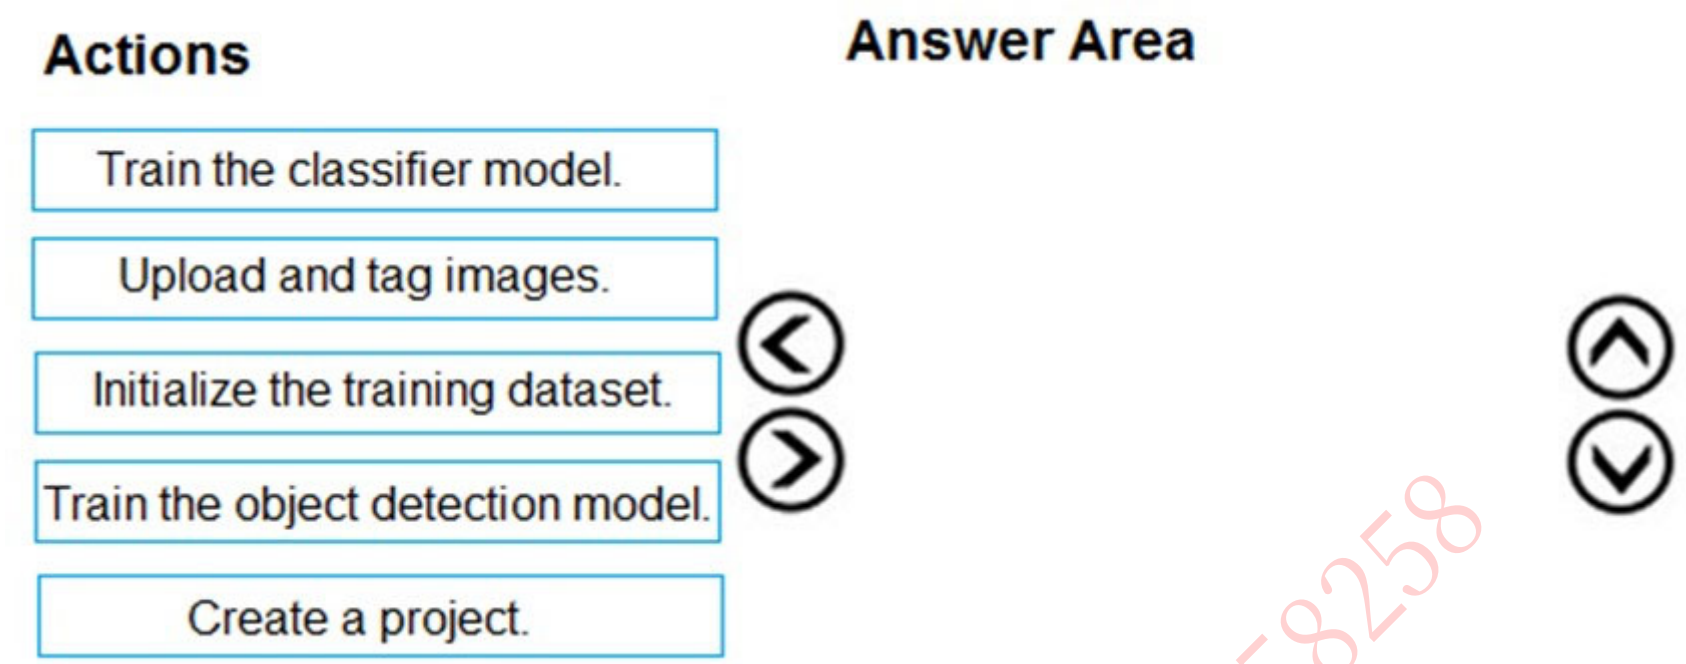

Correct Answer:

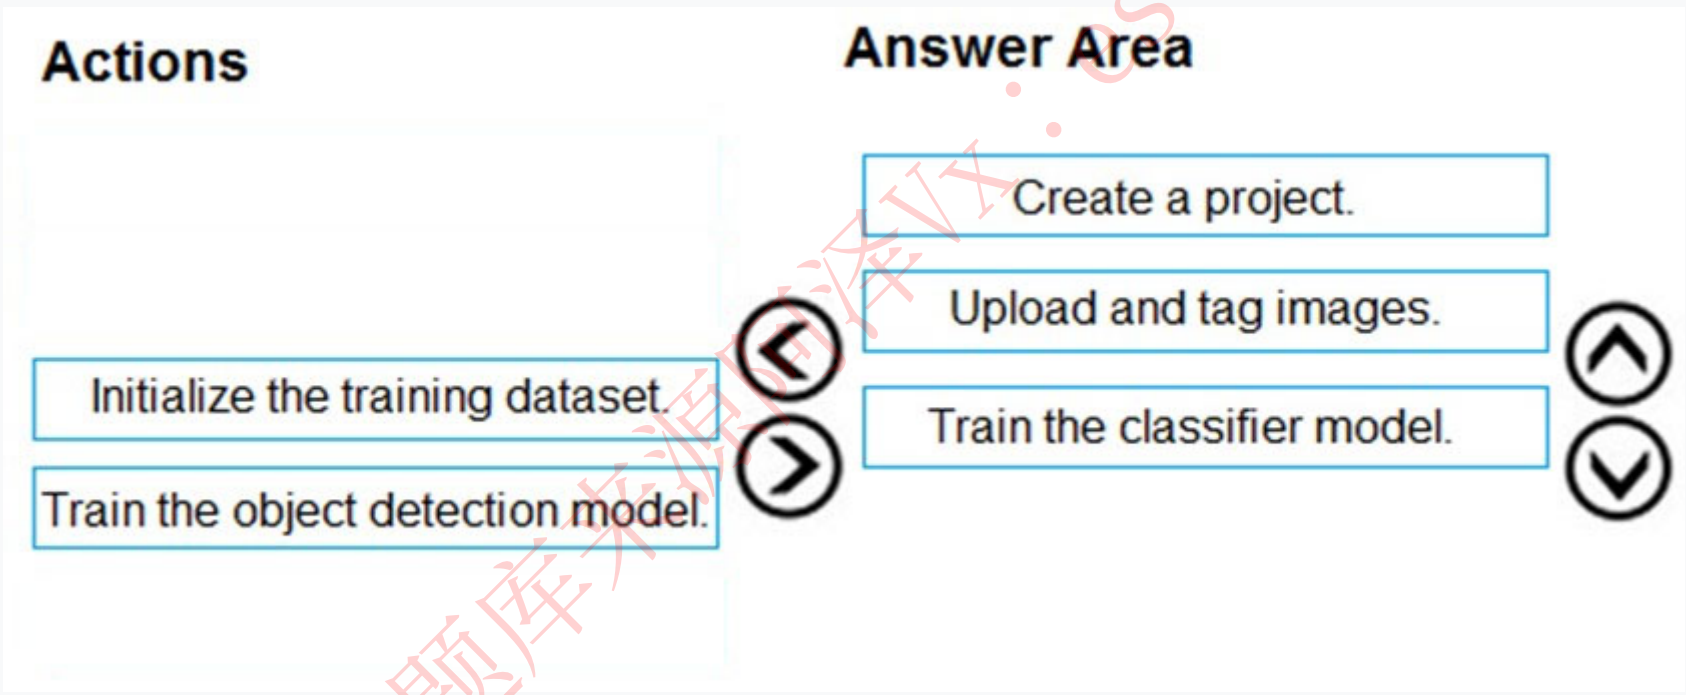

Step 1: Create a project -

Create a new project.

Step 2: Upload and tag the images

Choose training images. Then upload and tag the images.

Step 3: Train the classifier model.

Train the classifier -

Reference:

https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/getting-started-build-a-classifier

Question #9 *Topic 2*

HOTSPOT -

You are building a model that will be used in an iOS app.

You have images of cats and dogs. Each image contains either a cat or a dog.

You need to use the Custom Vision service to detect whether the images is of a cat or a dog.

How should you configure the project in the Custom Vision portal? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

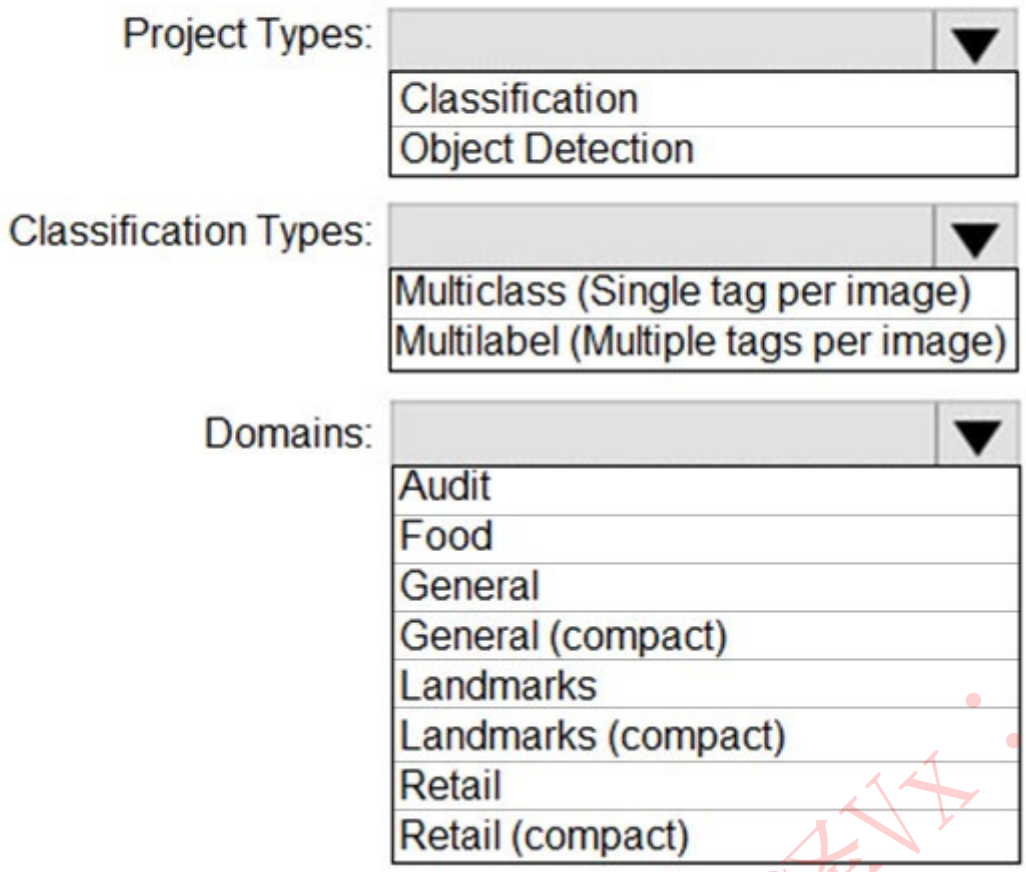

Correct Answer:

Answer Area

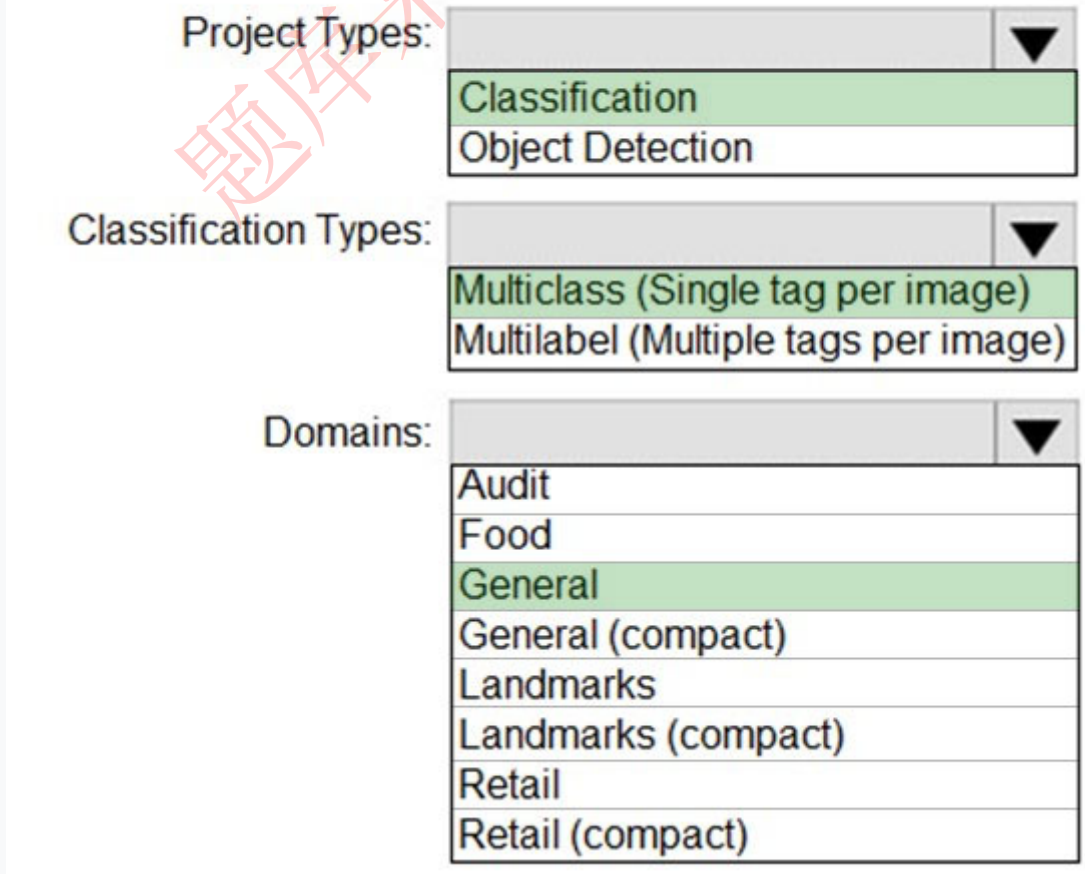

Box 1: Classification -

Incorrect Answers:

An object detection project is for detecting which objects, if any, from a set of candidates are present in an image.

Box 2: Multiclass -

A multiclass classification project is for classifying images into a set of tags, or target labels. An image can be assigned to one tag only.
Incorrect Answers:
A multilabel classification project is similar, but each image can have multiple tags assigned to it.

Box 3: General -
General: Optimized for a broad range of image classification tasks. If none of the other specific domains are appropriate, or if you're unsure of which domain to choose, select one of the General domains.
Reference:
https://cran.r-project.org/web/packages/AzureVision/vignettes/customvision.html

Question #10 *Topic 2*

You have an Azure Video Analyzer for Media (previously Video Indexer) service that is used to provide a search interface over company videos on your company's website.
You need to be able to search for videos based on who is present in the video.
What should you do?

- A. Create a person model and associate the model to the videos. **Most Voted**
- B. Create person objects and provide face images for each object.
- C. Invite the entire staff of the company to Video Indexer.
- D. Edit the faces in the videos.
- E. Upload names to a language model.

Correct Answer: *A*

Community vote distribution

A (94%) 6%

Question #11 *Topic 2*

You use the Custom Vision service to build a classifier.
After training is complete, you need to evaluate the classifier.
Which two metrics are available for review? Each correct answer presents a complete solution.
NOTE: Each correct selection is worth one point.

- A. recall **Most Voted**
- B. F-score
- C. weighted accuracy
- D. precision **Most Voted**
- E. area under the curve (AUC)

Correct Answer: *AD*

Community vote distribution

AD (100%)