Question #80 *Topic 3*

You have the following Python function.

```
def my_function(textAnalyticsClient, text):

    response = textAnalyticsClient.extract_key_phrases(documents = [text])[0]
    print("Key Phrases:")
    for phrase in response.key_phrases:
        print(phrase)
```

You call the function by using the following code.

my_function(text_analytics_client, "the quick brown fox jumps over the lazy dog")

Following 'Key phrases', what output will you receive?

A. The quick -
The lazy

B. jumps over the

C. quick brown fox
lazy dog

D. the quick brown fox jumps over the lazy dog

Correct Answer: *C*

Question #81 *Topic 3*

You have an Azure subscription.

You need to deploy an Azure AI Search resource that will recognize geographic locations.

Which built-in skill should you include in the skillset for the resource?

A. AzureOpenAIEmbeddingSkill

B. DocumentExtractionSkill

C. EntityRecognitionSkill

D. EntityLinkingSkill

Correct Answer: *C*

Topic 4 - Question Set 4

Question #1 *Topic 4*

HOTSPOT -

You are developing a text processing solution.

You develop the following method.

```
static void GetKeyPhrases(TextAnalyticsClient textAnalyticsClient, string text)
{
    var response = textAnalyticsClient.ExtractKeyPhrases(text);
    Console.WriteLine("Key phrases:");

    foreach (string keyphrase in response.Value)
    {
        Console.WriteLine($"\t{keyphrase}");
    }
}
```

You call the method by using the following code.

GetKeyPhrases(textAnalyticsClient, "the cat sat on the mat");

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
The call will output key phrases from the input string to the console.	○	○
The output will contain the following words: the, cat, sat, on, and mat.	○	○
The output will contain the confidence level for key phrases.	○	○

Correct Answer:

Answer Area

Statements	Yes	No
The call will output key phrases from the input string to the console.	○	○
The output will contain the following words: the, cat, sat, on, and mat.	○	○
The output will contain the confidence level for key phrases.	○	○

Box 1: Yes -

The Key Phrase Extraction API evaluates unstructured text, and for each JSON document, returns a list of key phrases.

Box 2: No -

'the' is not a key phrase.

This capability is useful if you need to quickly identify the main points in a collection of documents. For example, given input text "The food was delicious and there were wonderful staff", the service returns the main talking points: "food" and "wonderful staff".

Box 3: No -

Key phrase extraction does not have confidence levels.

Reference:

https://docs.microsoft.com/en-us/azure/cognitive-services/text-analytics/how-tos/text-analytics-how-to-keyword-extraction

Question #2 *Topic 4*

You deploy a web app that is used as a management portal for indexing in Azure Cognitive Search. The app is configured to use the primary admin key.
During a security review, you discover unauthorized changes to the search index. You suspect that the primary access key is compromised.
You need to prevent unauthorized access to the index management endpoint. The solution must minimize downtime.
What should you do next?

A. Regenerate the primary admin key, change the app to use the secondary admin key, and then regenerate the secondary admin key.

B. Change the app to use a query key, and then regenerate the primary admin key and the secondary admin key.

C. Regenerate the secondary admin key, change the app to use the secondary admin key, and then regenerate the primary key. **Most Voted**

D. Add a new query key, change the app to use the new query key, and then delete all the unused query keys.

Correct Answer: *C*

Community vote distribution

C (89%) 7%

Question #3 *Topic 4*

You have an existing Azure Cognitive Search service.
You have an Azure Blob storage account that contains millions of scanned documents stored as images and PDFs.
You need to make the scanned documents available to search as quickly as possible.
What should you do?

A. Split the data into multiple blob containers. Create a Cognitive Search service for each container. Within each indexer definition, schedule the same runtime execution pattern.

B. Split the data into multiple blob containers. Create an indexer for each container. Increase the search units. Within each indexer definition, schedule a sequential execution pattern.

C. Create a Cognitive Search service for each type of document.

D. Split the data into multiple virtual folders. Create an indexer for each folder. Increase the search units. Within each indexer definition, schedule the same runtime execution pattern. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (100%)

Question #4 *Topic 4*

You need to implement a table projection to generate a physical expression of an Azure Cognitive Search index.
Which three properties should you specify in the skillset definition JSON configuration table node? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. tableName **Most Voted**

B. generatedKeyName **Most Voted**

C. dataSource

D. dataSourceConnection

E. source **Most Voted**

Correct Answer: *ABE*

Community vote distribution

ABE (100%)

Question #5 *Topic 4*

HOTSPOT -

You are creating an enrichment pipeline that will use Azure Cognitive Search. The knowledge store contains unstructured JSON data and scanned PDF documents that contain text.

Which projection type should you use for each data type? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

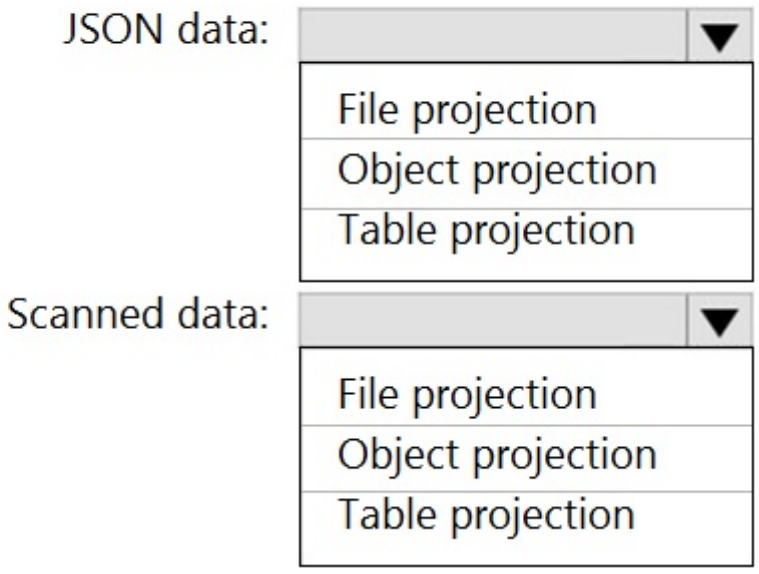

Correct Answer:

Answer Area

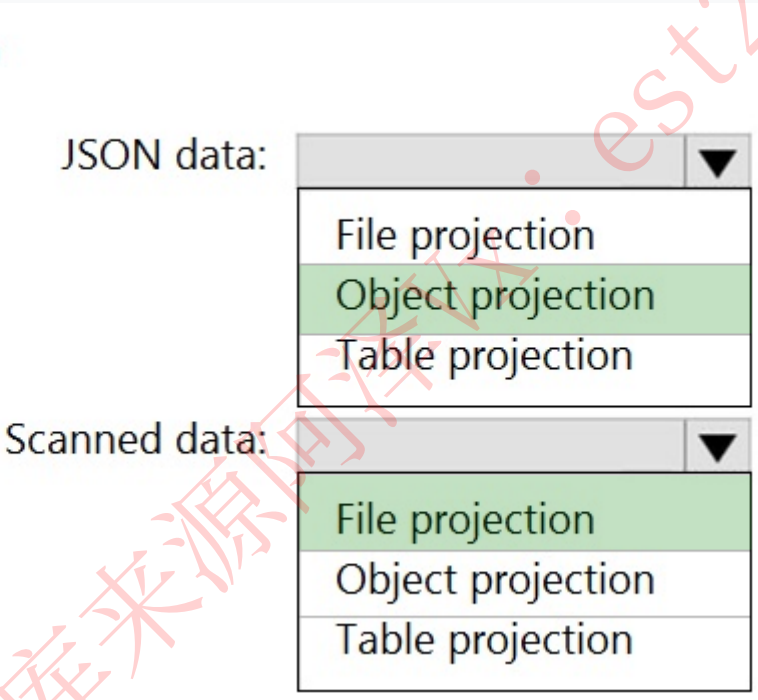

Box 1: Object projection -

Object projections are JSON representations of the enrichment tree that can be sourced from any node.

Box 2: File projection -

File projections are similar to object projections and only act on the normalized_images collection.

Reference:

https://docs.microsoft.com/en-us/azure/search/knowledge-store-projection-overview

Question #6 *Topic 4*

HOTSPOT -

You are building an Azure Cognitive Search custom skill.

You have the following custom skill schema definition.

```
{
    "@odata.type": "#Microsoft.Skills.Custom.WebApiSkill",
    "description": "My custom skill description",
    "uri": "https://contoso-webskill.azurewebsites.net/api/process",
    "context": "/document/organizations/*",
    "inputs": [
      {
        "name": "companyName",
        "source": "/document/organizations/*"
      }
    ],
    "outputs": [
      {
        "name": "companyDescription",
      }
    ]
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
`CompanyDescription` is available for indexing.	○	○
The definition calls a web API as part of the enrichment process.	○	○
The enrichment step is called only for the first organization under "`/document/organizations`".	○	○

Correct Answer:

Answer Area

Statements	Yes	No
`CompanyDescription` is available for indexing.	○	○
The definition calls a web API as part of the enrichment process.	○	○
The enrichment step is called only for the first organization under "`/document/organizations`".	○	○

Box 1: Yes -

Once you have defined a skillset, you must map the output fields of any skill that directly contributes values to a given field in your search index.

Box 2: Yes -

The definition is a custom skill that calls a web API as part of the enrichment process.

Box 3: No -

For each organization identified by entity recognition, this skill calls a web API to find the description of that organization.

Reference:

https://docs.microsoft.com/en-us/azure/search/cognitive-search-output-field-mapping

Question #7 *Topic 4*

You have the following data sources:

- Finance: On-premises Microsoft SQL Server database
- Sales: Azure Cosmos DB using the Core (SQL) API
- Logs: Azure Table storage

HR: Azure SQL database -

▪

You need to ensure that you can search all the data by using the Azure Cognitive Search REST API.

What should you do?

A. Configure multiple read replicas for the data in Sales.

B. Mirror Finance to an Azure SQL database. **Most Voted**

C. Ingest the data in Logs into Azure Data Explorer.

D. Ingest the data in Logs into Azure Sentinel.

Correct Answer: *B*

Community vote distribution

B (100%)

Question #8 *Topic 4*

You are developing a solution to generate a word cloud based on the reviews of a company's products.

Which Text Analytics REST API endpoint should you use?

A. keyPhrases **Most Voted**

B. sentiment

C. languages

D. entities/recognition/general

Correct Answer: *A*

Community vote distribution

A (95%) 5%

Question #9 *Topic 4*

DRAG DROP -

You have a web app that uses Azure Cognitive Search.

When reviewing billing for the app, you discover much higher than expected charges. You suspect that the query key is compromised.

You need to prevent unauthorized access to the search endpoint and ensure that users only have read only access to the documents collection. The solution must minimize app downtime.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

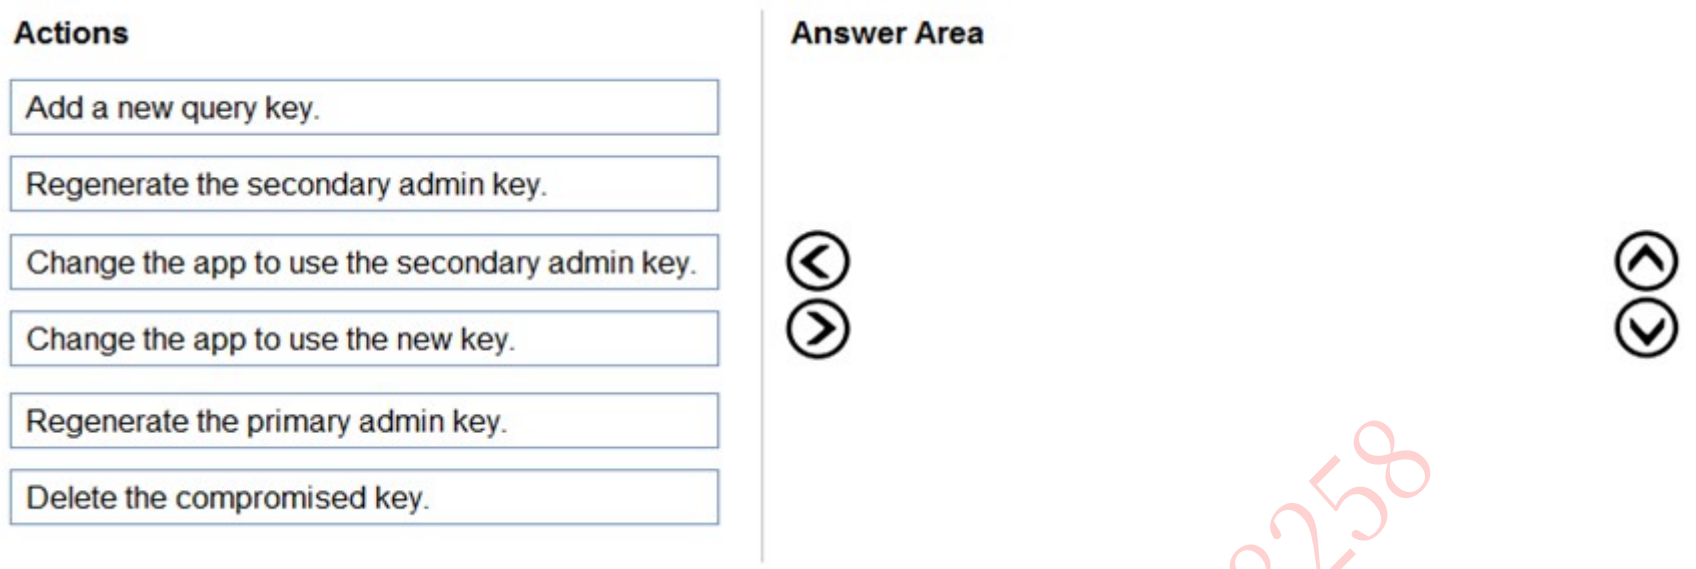

Correct Answer:

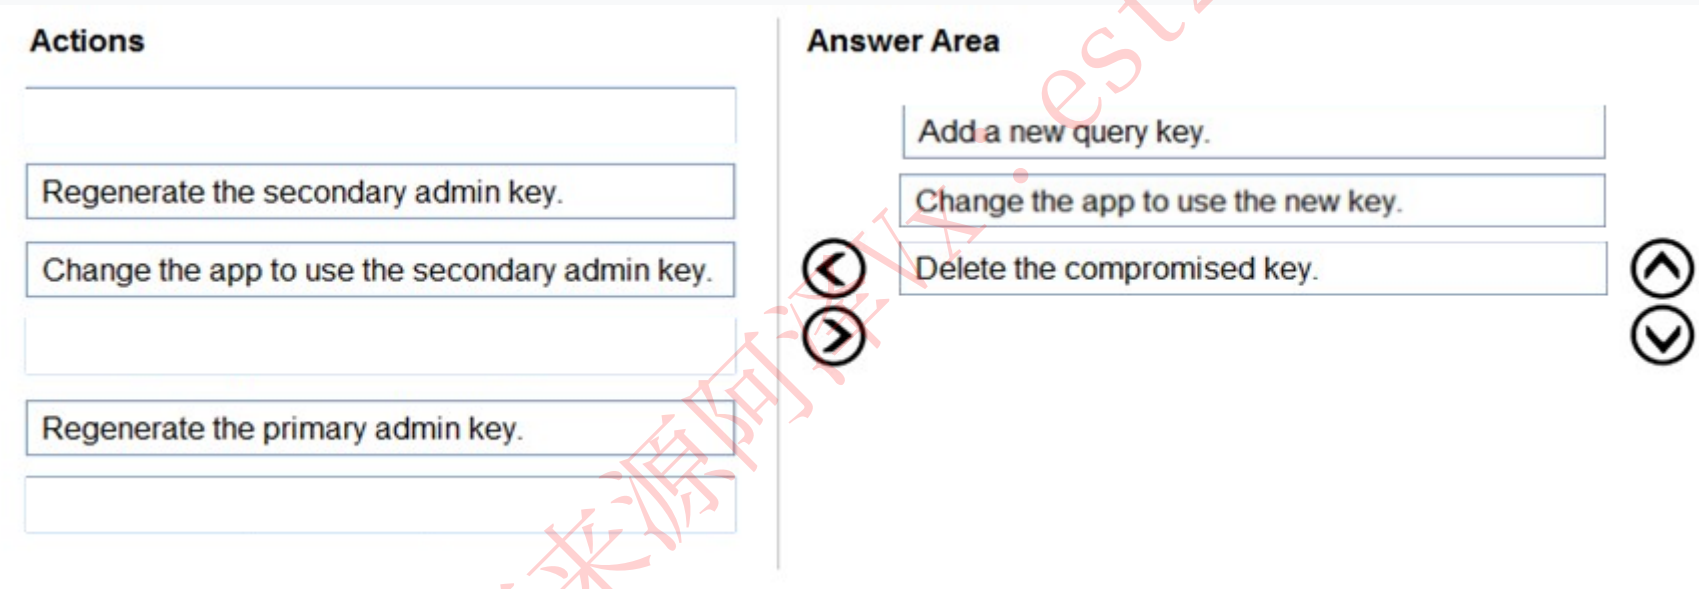

Reference:

https://docs.microsoft.com/en-us/azure/search/search-security-api-keys

Question #10 *Topic 4*

You are developing an application that will use Azure Cognitive Search for internal documents.
You need to implement document-level filtering for Azure Cognitive Search.
Which three actions should you include in the solution? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. Send Azure AD access tokens with the search request.

B. Retrieve all the groups.

C. Retrieve the group memberships of the user. **Most Voted**

D. Add allowed groups to each index entry. **Most Voted**

E. Create one index per group.

F. Supply the groups as a filter for the search requests. **Most Voted**

Correct Answer: *CDF*

Community vote distribution

Question #11 *Topic 4*

You have an Azure Cognitive Search solution and an enrichment pipeline that performs Sentiment Analysis on social media posts.
You need to define a knowledge store that will include the social media posts and the Sentiment Analysis results.
Which two fields should you include in the definition? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. storageContainer

B. storageConnectionString **Most Voted**

C. files

D. tables

E. objects **Most Voted**

Correct Answer: *BE*

Community vote distribution

Question #12 *Topic 4*

SIMULATION -
Use the following login credentials as needed:
To enter your username, place your cursor in the Sign in box and click on the username below.
To enter your password, place your cursor in the Enter password box and click on the password below.

Azure Username: admin@abc.com -

Azure Password: XXXXXXXXXXXX -
The following information is for technical support purposes only:

Lab Instance: 12345678 -

Task -
You need to create an Azure resource named solution12345678 that will index a sample database named realestate-us-sample. The solution must ensure that users can search the index in English for people, organizations, and locations.
To complete this task, sign in to the Azure portal.

Correct Answer: *See explanation below.*
Step 1 - Start the Import data wizard and create a data source
1. Sign in to the Azure portal with your Azure account.
2. Find your search service and on the Overview page, click Import data on the command bar to create and populate a search index.

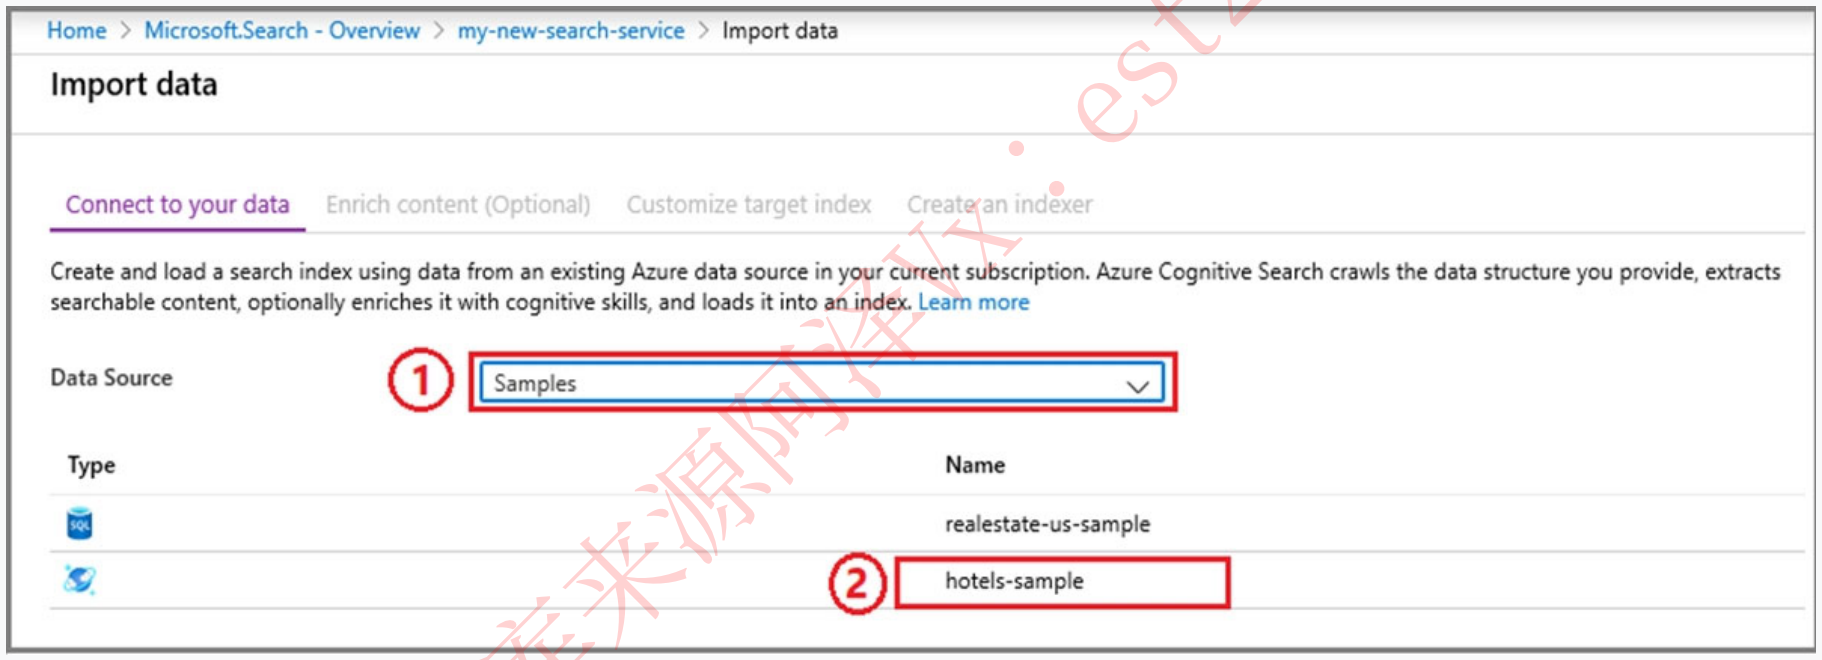

3. In the wizard, click Connect to your data, and select the sample database named realestate-us-sample
Step 2 - Skip the "Enrich content" page
The wizard supports the creation of an AI enrichment pipeline for incorporating the Cognitive Services AI algorithms into indexing.
We'll skip this step for now, and move directly on to Customize target index.

Step 3 - Configure index -
The solution must ensure that users can search the index in English for people, organizations, and locations.
Configure Searchable for the fields people, organizations, and locations.