

You build a bot.

You create an Azure Bot resource.

You need to deploy the bot to Azure.

What else should you create?

- A. only an app registration in Microsoft Azure Active Directory (Azure AD), part of Microsoft Entra, an Azure App Service instance, and an App Service plan **Most Voted**
- B. only an app registration in Microsoft Azure Active Directory (Azure AD), part of Microsoft Entra, an Azure Kubernetes Service (AKS) instance, and a container image
- C. only an Azure App Service instance, and an App Service plan
- D. only an Azure Machine Learning workspace and an app registration in Microsoft Azure Active Directory (Azure AD), part of Microsoft Entra

Correct Answer: A

Community vote distribution

A (63%)

C (37%)

You are building a chatbot by using the Microsoft Bot Framework SDK. The bot will be used to accept food orders from customers and allow the customers to customize each food item.

You need to configure the bot to ask the user for additional input based on the type of item ordered. The solution must minimize development effort.

Which two types of dialogs should you use? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. adaptive
- B. action
- C. waterfall **Most Voted**
- D. prompt **Most Voted**
- E. input

Correct Answer: CD

Community vote distribution

CD (100%)

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You are building a chatbot that will use question answering in Azure Cognitive Service for Language.

You have a PDF named Doc1.pdf that contains a product catalogue and a price list.

You upload Doc1.pdf and train the model.

During testing, users report that the chatbot responds correctly to the following question: What is the price of ?

The chatbot fails to respond to the following question: How much does cost?

You need to ensure that the chatbot responds correctly to both questions.

Solution: From Language Studio, you create an entity for cost, and then retrain and republish the model.

Does this meet the goal?

A. Yes

B. No **Most Voted**

Correct Answer: B

Community vote distribution

B (100%)

DRAG DROP

-

You have a monitoring solution that uses the Azure AI Anomaly Detector service.

You provision a server named Server1 that has intermittent internet access.

You need to deploy the Azure AI Anomaly Detector to Server1.

Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

Query the prediction endpoint on Server1.

From Server1, run the docker push command.

Install the Docker Engine on Server1.

Query the prediction endpoint of the Azure AI Anomaly Detector in Azure.

From Server1, run the docker run command.

From Server1, run the docker pull command.

Answer Area

1

2

3

4

>

<

⬆

⬇

Correct Answer:

Answer Area

1

Install the Docker Engine on Server1.

2

From Server1, run the docker pull command.

3

From Server1, run the docker run command.

4

Query the prediction endpoint on Server1.

You have an Azure subscription. The subscription contains an Azure OpenAI resource that hosts a GPT-4 model named Model1 and an app named App1. App1 uses Model1.

You need to ensure that App1 will NOT return answers that include hate speech.

What should you configure for Model1?

- A. the Frequency penalty parameter
- B. abuse monitoring
- C. a content filter **Most Voted**
- D. the Temperature parameter

Correct Answer: C

Community vote distribution

C (100%)

You have an Azure subscription. The subscription contains an Azure OpenAI resource that hosts a GPT-3.5 Turbo model named Model1.

You configure Model1 to use the following system message: "You are an AI assistant that helps people solve mathematical puzzles. Explain your answers as if the request is by a 4-year-old."

Which type of prompt engineering technique is this an example of?

- A. few-shot learning
- B. affordance
- C. chain of thought
- D. priming **Most Voted**

Correct Answer: D

Community vote distribution

D (77%)C (18%)5%

HOTSPOT

You build a chatbot by using Azure OpenAI Studio.

You need to ensure that the responses are more deterministic and less creative.

Which two parameters should you configure? To answer, select the appropriate parameters in the answer area.

NOTE: Each correct answer is worth one point.

Answer Area

Chat session

Clear chat

View code

Show raw JSON

Start chatting

Test your assistant by sending queries below. Then adjust your assistant setup to improve the assistant's responses.

Type user query here. (Shift + Enter for new line)

Configuration

Deployment

Parameters

Max response

800

Temperature

0.7

Top P

0.9

Stop sequence

Stop sequences

Frequency penalty

0

Presence penalty

0

Learn more

Current token count

Input tokens progress indicator

1/4000

Answer Area

Chat session

Clear chat View code Show raw JSON



Start chatting

Test your assistant by sending queries below. Then

Configuration

Deployment Parameters

Max response 800

Temperature 0.7

Top P 0.9

Correct Answer:

Question #4

Topic 7

You are building a chatbot for a travel agent. The chatbot will use the Azure OpenAI GPT 3.5 model and will be used to make travel reservations.

You need to maximize the accuracy of the responses from the chatbot.

What should you do?

- A. Configure the model to include data from the travel agent's database. **Most Voted**
- B. Set the Top P parameter for the model to 0.
- C. Set the Temperature parameter for the model to 0.
- D. Modify the system message used by the model to specify that the answers must be accurate.

Correct Answer: A

Community vote distribution

A (69%)

D (24%)

3%

Question #5

Topic 7

You build a chatbot that uses the Azure OpenAI GPT 3.5 model.

You need to improve the quality of the responses from the chatbot. The solution must minimize development effort.

What are two ways to achieve the goal? Each correct answer presents a complete solution.

NOTE: Each correct answer is worth one point.

- A. Fine-tune the model.
- B. Provide grounding content. **Most Voted**
- C. Add sample request/response pairs. **Most Voted**
- D. Retrain the language model by using your own data.
- E. Train a custom large language model (LLM).

Correct Answer: BC

Community vote distribution

BC (100%)

HOTSPOT

-

You have an Azure subscription that contains an Azure OpenAI resource named AI1.

You build a chatbot that will use AI1 to provide generative answers to specific questions.

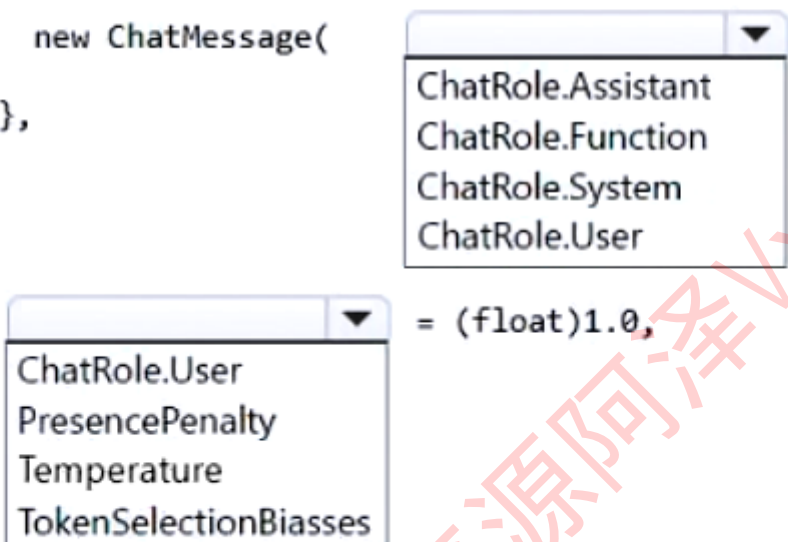
You need to ensure that the responses are more creative and less deterministic.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
new ChatCompletionsOptions()  
{  
    Messages =  
    {  
        new ChatMessage(  
            , @""),  
    },  
    = (float)1.0,  
    MaxTokens = 800,  
});
```



Answer Area

```
new ChatCompletionOptions()
```

Question #7

Topic 7

DRAG DROP

You have an Azure subscription that contains an Azure OpenAI resource named AI1.

You plan to build an app named App1 that will write press releases by using AI1.

You need to deploy an Azure OpenAI model for App1. The solution must minimize development effort.

Which three actions should you perform in sequence in Azure OpenAI Studio? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

- Create a deployment that uses the text-embedding-ada-002 model.
- Apply the Marketing Writing Assistant system message template.
- Apply the Default system message template.
- Create a deployment that uses the GPT-35 Turbo model.
- Deploy the solution to a new web app.

Answer Area

- 1
- 2
- 3

Correct Answer:

Answer Area

1

Create a deployment that uses the GPT-35 Turbo model.

2

Apply the Marketing Writing Assistant system message template.

3

Deploy the solution to a new web app.

HOTSPOT

-

You have an Azure subscription that contains an Azure OpenAI resource named AI1.

You build a chatbot that will use AI1 to provide generative answers to specific questions.

You need to ensure that the responses are more creative and less deterministic.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
response = openai.ChatCompletion.create(  
    engine="dgw-aoai-gpt35",  
    messages = [{"role": , "content":""}],  
     =1,  
    max_tokens=800,  
    stop=None)
```

Options for role: "assistant", "function", "system", "user"

Options for temperature: Frequency_penalty, Presence_penalty, temperature, token_selection_biasses

Correct Answer:**Answer Area**

```
response = openai.ChatCompletion.create(  
    engine="dgw-aoai-gpt35",  
    messages = [{"role": , "content":""}],  
     =1,  
    max_tokens=800,  
    stop=None)
```

Options for role: "assistant", "function", "system", "user"

Options for temperature: Frequency_penalty, Presence_penalty, temperature, token_selection_biasses

题库来源阿泽Vx: est258258