

You have an Azure subscription that contains an Azure OpenAI resource named AI1 and a user named User1.

You need to ensure that User1 can perform the following actions in Azure OpenAI Studio:

- Identify resource endpoints.
- View models that are available for deployment.
- Generate text and images by using the deployed models.

The solution must follow the principle of least privilege.

Which role should you assign to User1?

- A. Cognitive Services OpenAI User

B. Cognitive Services Contributor

C. Contributor

D. Cognitive Services OpenAI Contributor

Correct Answer: A

Community vote distribution

A (100%)

## HOTSPOT

You have a chatbot that uses Azure OpenAI to generate responses.

You need to upload company data by using Chat playground. The solution must ensure that the chatbot uses the data to answer user questions.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

## Answer Area

```
completion = openai.  
    .create(  
    messages=[{"role": "user", "content": "What are the differences between Azure Machine Learning and Azure AI services?"}],  
    deployment_id=os.environ.get("AOAIDeploymentId"),  
    dataSources=[  
        {  
            "type":  
            "parameters": {  
                "endpoint": os.environ.get("SearchEndpoint"),  
                "key": os.environ.get("SearchKey"),  
                "indexName": os.environ.get("SearchIndex"),  
                ...  
            }  
        }  
    ]  
)
```

## Correct Answer:

```
Answer Area  
completion = openai.  
    ChatCompletion  
    .create(  
    messages=[{"role": "user", "content": "What are the differences between Azure Machine Learning and Azure AI services?"}],  
    deployment_id=os.environ.get("AOAIDeploymentId"),  
    dataSources=[  
        {  
            "type":  
            AzureCognitiveSearch  
            "parameters": {  
                "endpoint": os.environ.get("SearchEndpoint"),  
                "key": os.environ.get("SearchKey"),  
                "indexName": os.environ.get("SearchIndex"),  
                ...  
            }  
        }  
    ]  
)
```

## HOTSPOT

-

You are building an app that will provide users with definitions of common AI terms.

You create the following Python code.

```
...
openai.api_key = key
openai.api_base = endpoint
response = openai.ChatCompletion.create(
    engine=deployment_name
    messages=[
        {"role": "system", "content": "You are a helpful assistant."},
        {"role": "user", "content": "What is an LLM?"}
    ]
)

print(response['choices'][0]['message']['content'])
...
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

**Answer Area****Statements****Yes****No**

The response will contain an explanation of large language models (LLMs) that has a high degree of certainty.

☐☐

Changing "What is an LLM?" to "What is an LLM in the context of AI models?" will produce the intended response.

☐☐

Changing "You are a helpful assistant." to "You must answer only within the context of AI language models." will produce the intended response.

☐☐**Correct Answer:****Answer Area****Statements****Yes****No**

The response will contain an explanation of large language models (LLMs) that has a high degree of certainty.

☐☒

Changing "What is an LLM?" to "What is an LLM in the context of AI models?" will produce the intended response.

☒☐

Changing "You are a helpful assistant." to "You must answer only within the context of AI language models." will produce the intended response.

☒☐

You have an Azure subscription that contains an Azure OpenAI resource named AI1.

You build a chatbot that uses AI1 to provide generative answers to specific questions.

You need to ensure that the chatbot checks all input and output for objectionable content.

Which type of resource should you create first?

- A. Microsoft Defender Threat Intelligence (Defender TI)
- B. Azure AI Content Safety
- C. Log Analytics
- D. Azure Machine Learning

**Correct Answer:** B

*Community vote distribution*

B (100%)

You have an Azure subscription that contains an Azure AI Content Safety resource named CS1.

You create a test image that contains a circle.

You submit the test image to CS1 by using the curl command and the following command-line parameters.

```
--data-raw '{
  "image": {
    "content": "<base_64_string>"
  },
  "categories": [
    "Violence"
  ],
  "outputType": "EightSeverityLevels"
}
```

- What should you expect as the output?
- A. 0
  - B. 0.0
  - C. 7
  - D. 100

**Correct Answer:** C

*Community vote distribution*

A (100%)

You have an Azure subscription.

You are building a social media app that will enable users to share images.

You need to ensure that inappropriate content uploaded by the users is blocked. The solution must minimize development effort.

What are two tools that you can use? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. Azure AI Document Intelligence
- B. Microsoft Defender for Cloud Apps
- C. Azure AI Content Safety Most Voted
- D. Azure AI Vision
- E. Azure AI Custom Vision

**Correct Answer:** CE

*Community vote distribution*

CD (50%)

C (50%)

题库来源阿泽Vx: est258258

## HOTSPOT

-

You have an Azure subscription that contains an Azure AI Content Safety resource named CS1.

You need to call CS1 to identify whether a user request contains hateful language.

How should you complete the command? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

## Answer Area

```
curl --location --request POST 'https://cs1.cognitiveservices.azure.com/
--header 'Ocp-Apim-Subscription-Key: <your_subscription_key>' \
--header 'Content-Type: application/json' \
--data-raw '{
  "text": "What is the weather forecast for Seattle",
  "categories": ["Hate"]
  "blocklistNames": [
    "string"
  ],
  "haltOnBlocklistHit": true,
  "outputType": "FourSeverityLevels"
}'
```

completions/  
contentsafety/  
healthinsights/  
language/

completions  
embeddings  
text:analyze  
text/blocklists

?api-version=2023-10-01' \

## Correct Answer:

```
Answer Area
curl --location --request POST 'https://cs1.cognitiveservices.azure.com/
--header 'Ocp-Apim-Subscription-Key: <your_subscription_key>' \
--header 'Content-Type: application/json' \
--data-raw '{
  "text": "What is the weather forecast for Seattle",
  "categories": ["Hate"]
  "blocklistNames": [
    "string"
  ],
  "haltOnBlocklistHit": true,
  "outputType": "FourSeverityLevels"
}'
```

completions/  
contentsafety/  
healthinsights/  
language/

completions  
embeddings  
text:analyze  
text/blocklists

?api-version=2023-10-01' \

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure subscription that contains an Azure OpenAI resource named AI1 and an Azure AI Content Safety resource named CS1.

You build a chatbot that uses AI1 to provide generative answers to specific questions and CS1 to check input and output for objectionable content.

You need to optimize the content filter configurations by running tests on sample questions.

Solution: From Content Safety Studio, you use the Protected material detection feature to run the tests.

Does this meet the requirement?

- A. Yes
- B. No

**Correct Answer:** B

*Community vote distribution*

A (100%)

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure subscription that contains an Azure OpenAI resource named AI1 and an Azure AI Content Safety resource named CS1.

You build a chatbot that uses AI1 to provide generative answers to specific questions and CS1 to check input and output for objectionable content.

You need to optimize the content filter configurations by running tests on sample questions.

Solution: From Content Safety Studio, you use the Moderate text content feature to run the tests.

Does this meet the requirement?

- A. Yes
- B. No

**Correct Answer:** A

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure subscription that contains an Azure OpenAI resource named AI1 and an Azure AI Content Safety resource named CS1.

You build a chatbot that uses AI1 to provide generative answers to specific questions and CS1 to check input and output for objectionable content.

You need to optimize the content filter configurations by running tests on sample questions.

Solution: From Content Safety Studio, you use the Monitor online activity feature to run the tests.

Does this meet the requirement?

- A. Yes
- B. No

Correct Answer: B

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure subscription that contains an Azure OpenAI resource named AI1 and an Azure AI Content Safety resource named CS1.

You build a chatbot that uses AI1 to provide generative answers to specific questions and CS1 to check input and output for objectionable content.

You need to optimize the content filter configurations by running tests on sample questions.

Solution: From Content Safety Studio, you use the Safety metaprompt feature to run the tests.

Does this meet the requirement?

- A. Yes
- B. No

Correct Answer: B

Community vote distribution

B (100%)



## HOTSPOT

-

You have an Azure subscription that contains an Azure AI Content Safety resource.

You are building a social media app that will enable users to share images.

You need to configure the app to moderate inappropriate content uploaded by the users.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

**Answer Area**

```
public static void Analyze(AnalyzeImageOptions request)
{
    var endpoint = Environment.GetEnvironmentVariable("ENDPOINT");
    var key = Environment.GetEnvironmentVariable("KEY");
    var client = new 
        AnalyzeTextOptions
        BlocklistClient
        ContentSafetyClient
        TextCategoriesAnalysis
     (new Uri(endpoint), new AzureKeyCredential(key));

    return 
        AnalyzeImage(request)
        client.AnalyzeImage(request)
        client.AnalyzeText(request)
        request.AnalyzeImage(client)
     (request);
}
```

**Correct Answer:**

**Answer Area**

```
public static void Analyze(AnalyzeImageOptions request)
{
    var endpoint = Environment.GetEnvironmentVariable("ENDPOINT");
    var key = Environment.GetEnvironmentVariable("KEY");
    var client = new 
        AnalyzeTextOptions
        BlocklistClient
        ContentSafetyClient
        TextCategoriesAnalysis
     (new Uri(endpoint), new AzureKeyCredential(key));

    return 
        AnalyzeImage(request)
        client.AnalyzeImage(request)
        client.AnalyzeText(request)
        request.AnalyzeImage(client)
     (request);
}
```

You have an Azure subscription that contains an Azure AI Content Safety resource named CS1.

You plan to build an app that will analyze user-generated documents and identify obscure offensive terms.

You need to create a dictionary that will contain the offensive terms. The solution must minimize development effort.

What should you use?

- A. a text classifier
- B. language detection
- C. text moderation
- D. a blocklist

**Correct Answer:** D

*Community vote distribution*

D (100%)

题库来源阿泽Vx: est258258